

DEVELOPER-FRIENDLY TASKQUEUES

WHAT WE LEARNED BUILDING MRQ

& WHAT YOU SHOULD ASK YOURSELF BEFORE CHOOSING ONE

A photograph of a rocket launching from a launch pad. The rocket is white with 'SPACEX' printed vertically on its side. It is lifting off from a complex metal structure, likely a mobile service tower. A massive plume of white smoke and orange fire erupts from the base of the rocket, partially obscuring the launch pad. The background is a clear blue sky.

Sylvain Zimmer / @sylvinus

PyParis 2017

/usr/bin/whoami

- ▶ (SpaceX nerd)
- ▶ Founder, dotConferences
- ▶ CTO Pricing Assistant
- ▶ Co-organizer Paris.py meetup
- ▶ User of Python taskqueues for 10+ years
- ▶ Main contributor of MRQ

4 years ago...

2 people clipped this slide

A Python Task Queue Story

Why and how we migrated from Celery to RQ

Paris.py #2
22-07-2013

Sylvain Zimmer
@sylvinus

1 of 28

Edit Privacy Settings Analytics FREE

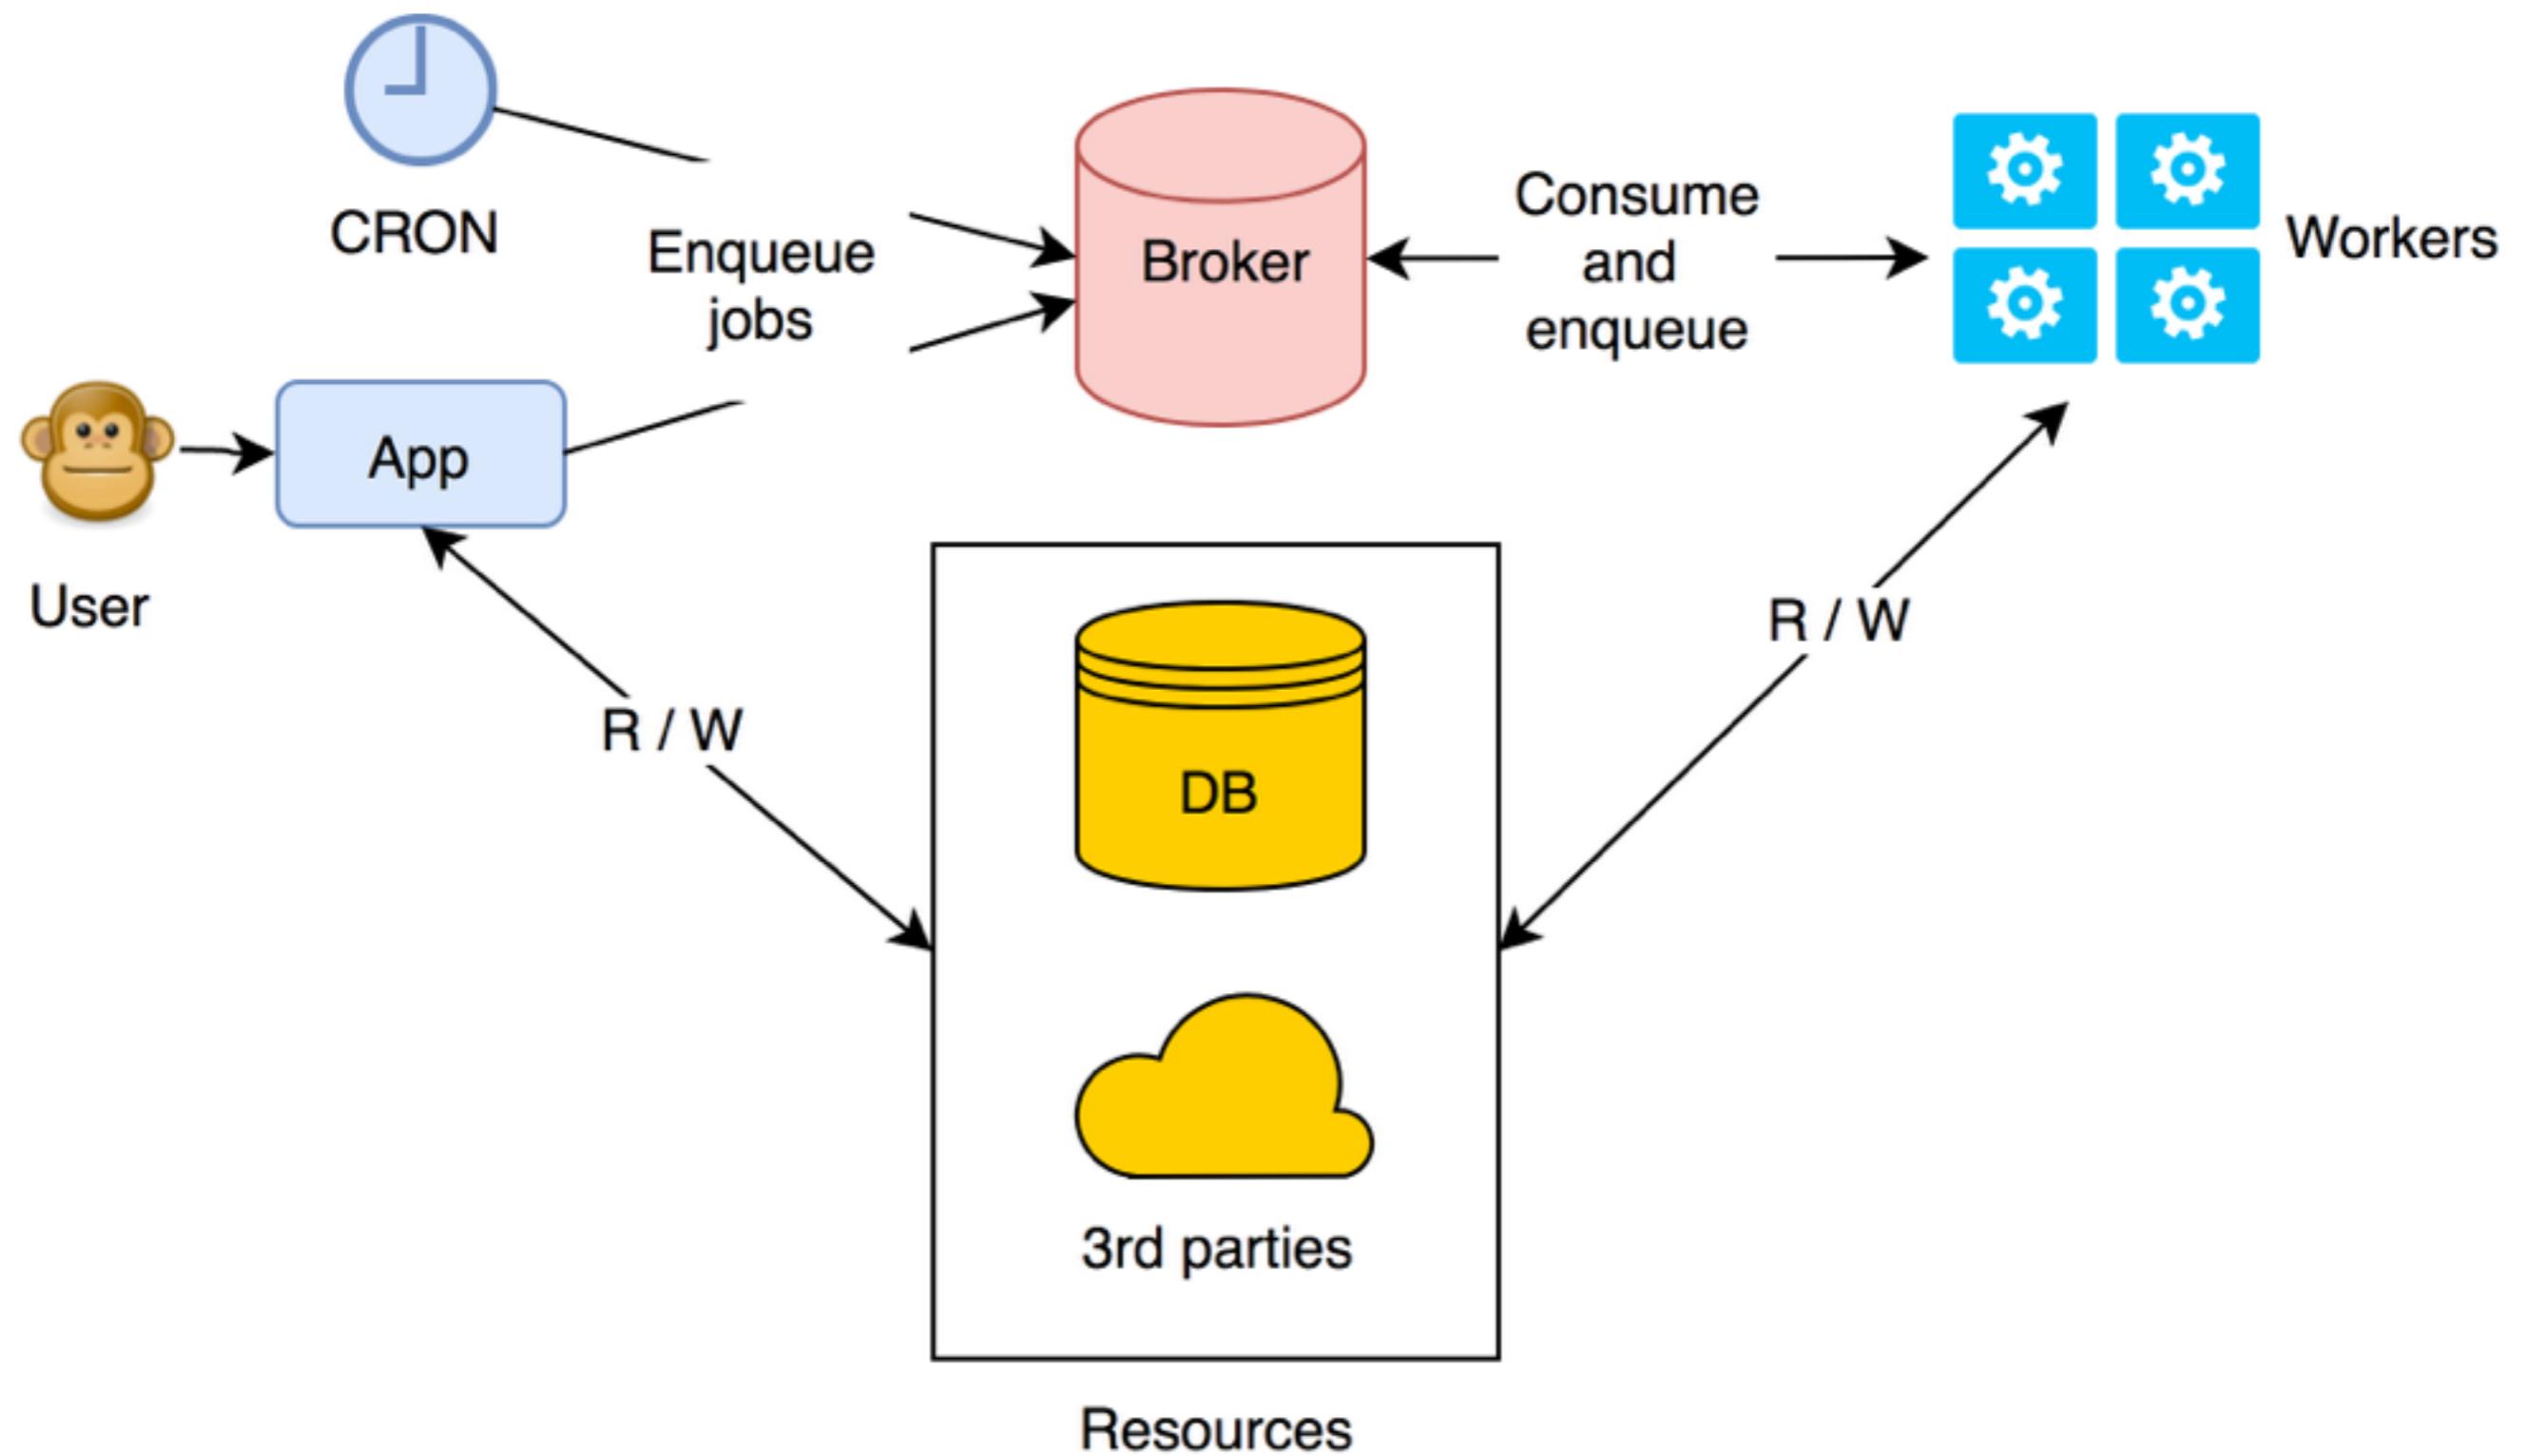
13,717 views

OMG !1!!!

Why and how Pricing Assistant migrated from Celery to RQ - Paris.py #2

TASKQUEUE FUNDAMENTALS





Credit: Adrien Di Pasquale

A typical job/task

```
def send_an_email(email_type, user):  
    html = template(email_type, user)  
    status = email.send(html, user["email"])  
KERNEL PANIC  
    metrics.send("email_%s" % status, 1)  
  
    return status
```

Task properties

Re-entrant < Idempotent < Nullipotent

- ▶ Safe to interrupt and then retry
- ▶ Safe to call multiple times
- ▶ Result will be the same
- ▶ Free of side-effects

```
def reentrant(a):  
    value = a + random()  
  
    db.insert(value)
```

```
def idempotent(key, value):  
    db.update(key, value)
```

```
def nullipotent(a):  
    return a ** 2
```

Other task properties & best practices

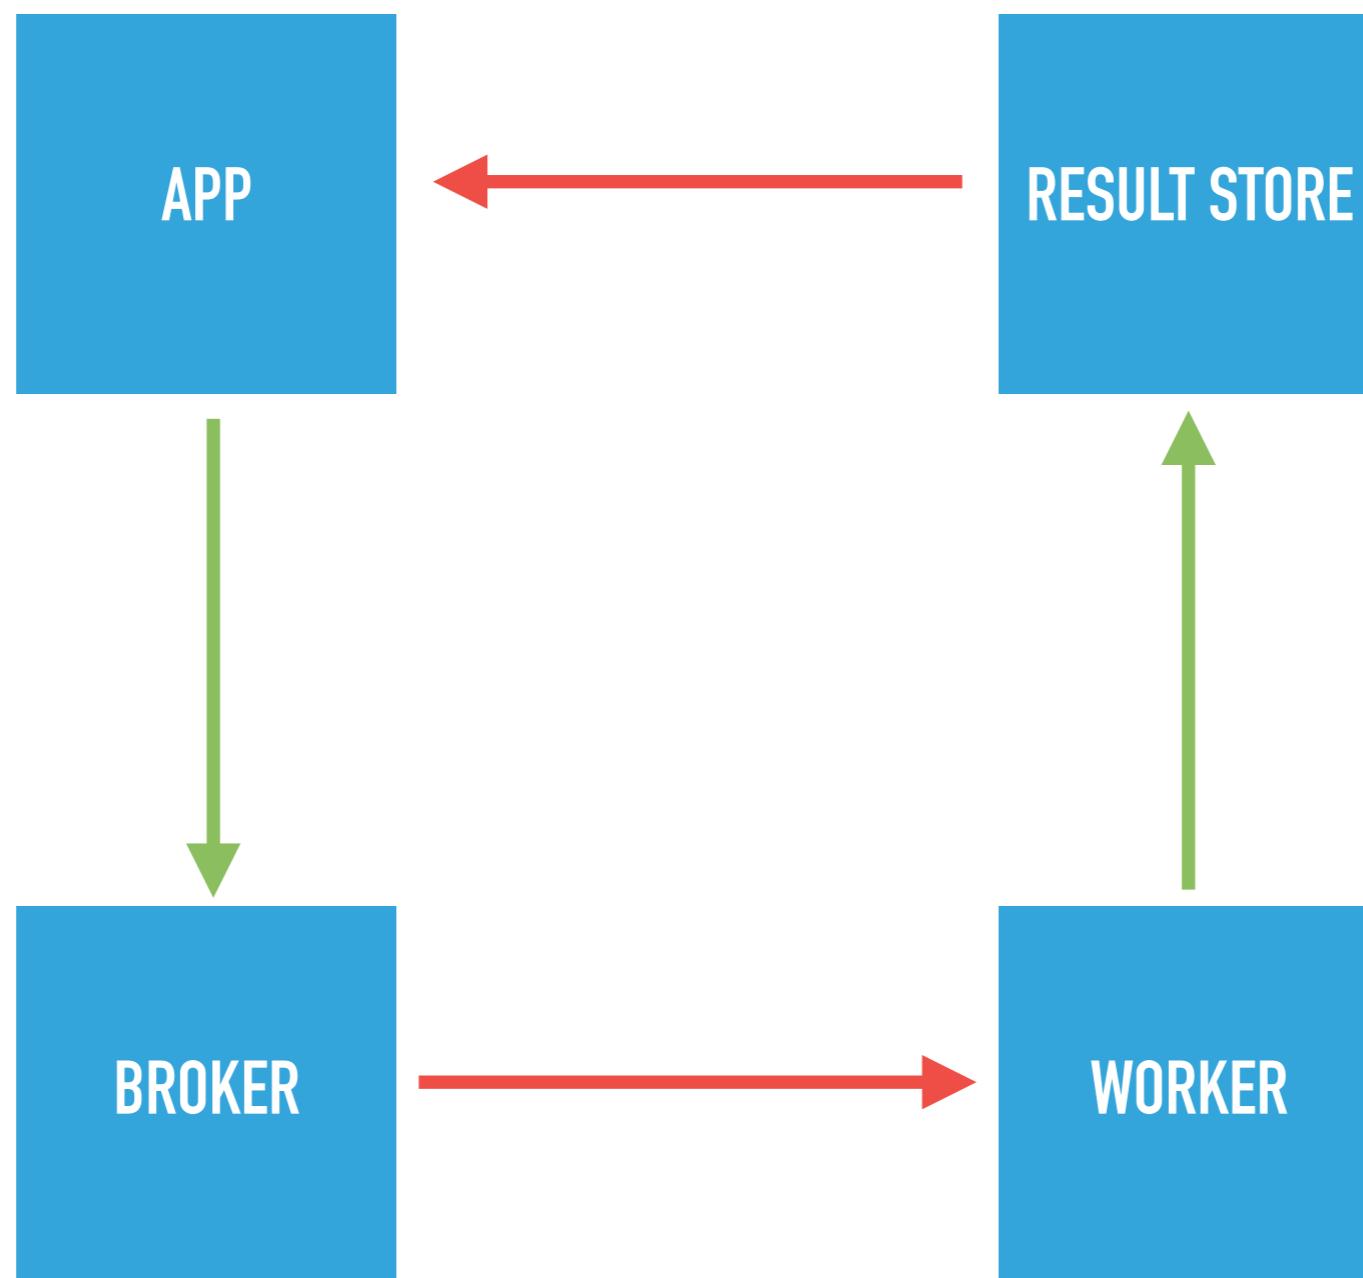
- ▶ Serializable args, serializable result
- ▶ Args validation / documentation
- ▶ Least args possible
- ▶ Canonical path vs. registration
- ▶ Concurrent safety
- ▶ Statuses

Coroutines vs. Threads vs. Processes

- ▶ IO-bound tasks vs. CPU-bound tasks
- ▶ Threads offer few benefits for a Python worker (GIL)
- ▶ Coroutines/Greenlets are ideal for IO-bound tasks
- ▶ Processes are required for CPU-bound tasks
- ▶ If you have heterogenous tasks, your TQ should support both!

```
$ mrq-worker --greenlets 25 --processes 4
```

Performance: latency & throughput



OPS & TOOLING





MURPHY'S LAW DOESN'T MEAN
THAT SOMETHING BAD WILL HAPPEN.
IT MEANS THAT WHATEVER CAN HAPPEN, WILL HAPPEN.

Errors

- ▶ Exception handlers
- ▶ Timeouts
- ▶ Retry rules
- ▶ Sentry & friends
- ▶ gevent: test your tracebacks!
- ▶ Priorities
- ▶ Human process to manage failed tasks!

Task visibility

- ▶ Tasks by status, path, worker, ...
- ▶ Tracebacks & current stack
- ▶ Logs
- ▶ Timing info
- ▶ Cancel / Kill / Move tasks
- ▶ Progress

tasks.exports.export.One	<pre>{\n "format": "csv",\n "skip_delete": false,\n "dailytime": 8,\n "static_upload": true,\n "store": "demo2",\n "remote_filename": "demo2.csv"\n}</pre>	started	queued 2 hours ago	exports	5936532c0009d5000 bfbc378	Logs	Result
593e4a6e6636e9000be35947		85%	started 2 hours ago			Command	Requeue

Memory leaks

- ▶ Workers = long-running processes
- ▶ gevent makes debugging harder
- ▶ Watch out for global variables or mutable class attributes!
- ▶ Python's ecosystem is surprisingly poor in this area
- ▶ guppy, objgraph can usually help

```
$ pip install guppy
$ python
>>> from mrq.context import setup_context, run_task
>>> setup_context()
>>> from guppy import hpy
>>> hp = hpy()
```

Then, wrap your memory-intensive task or code around guppy calls

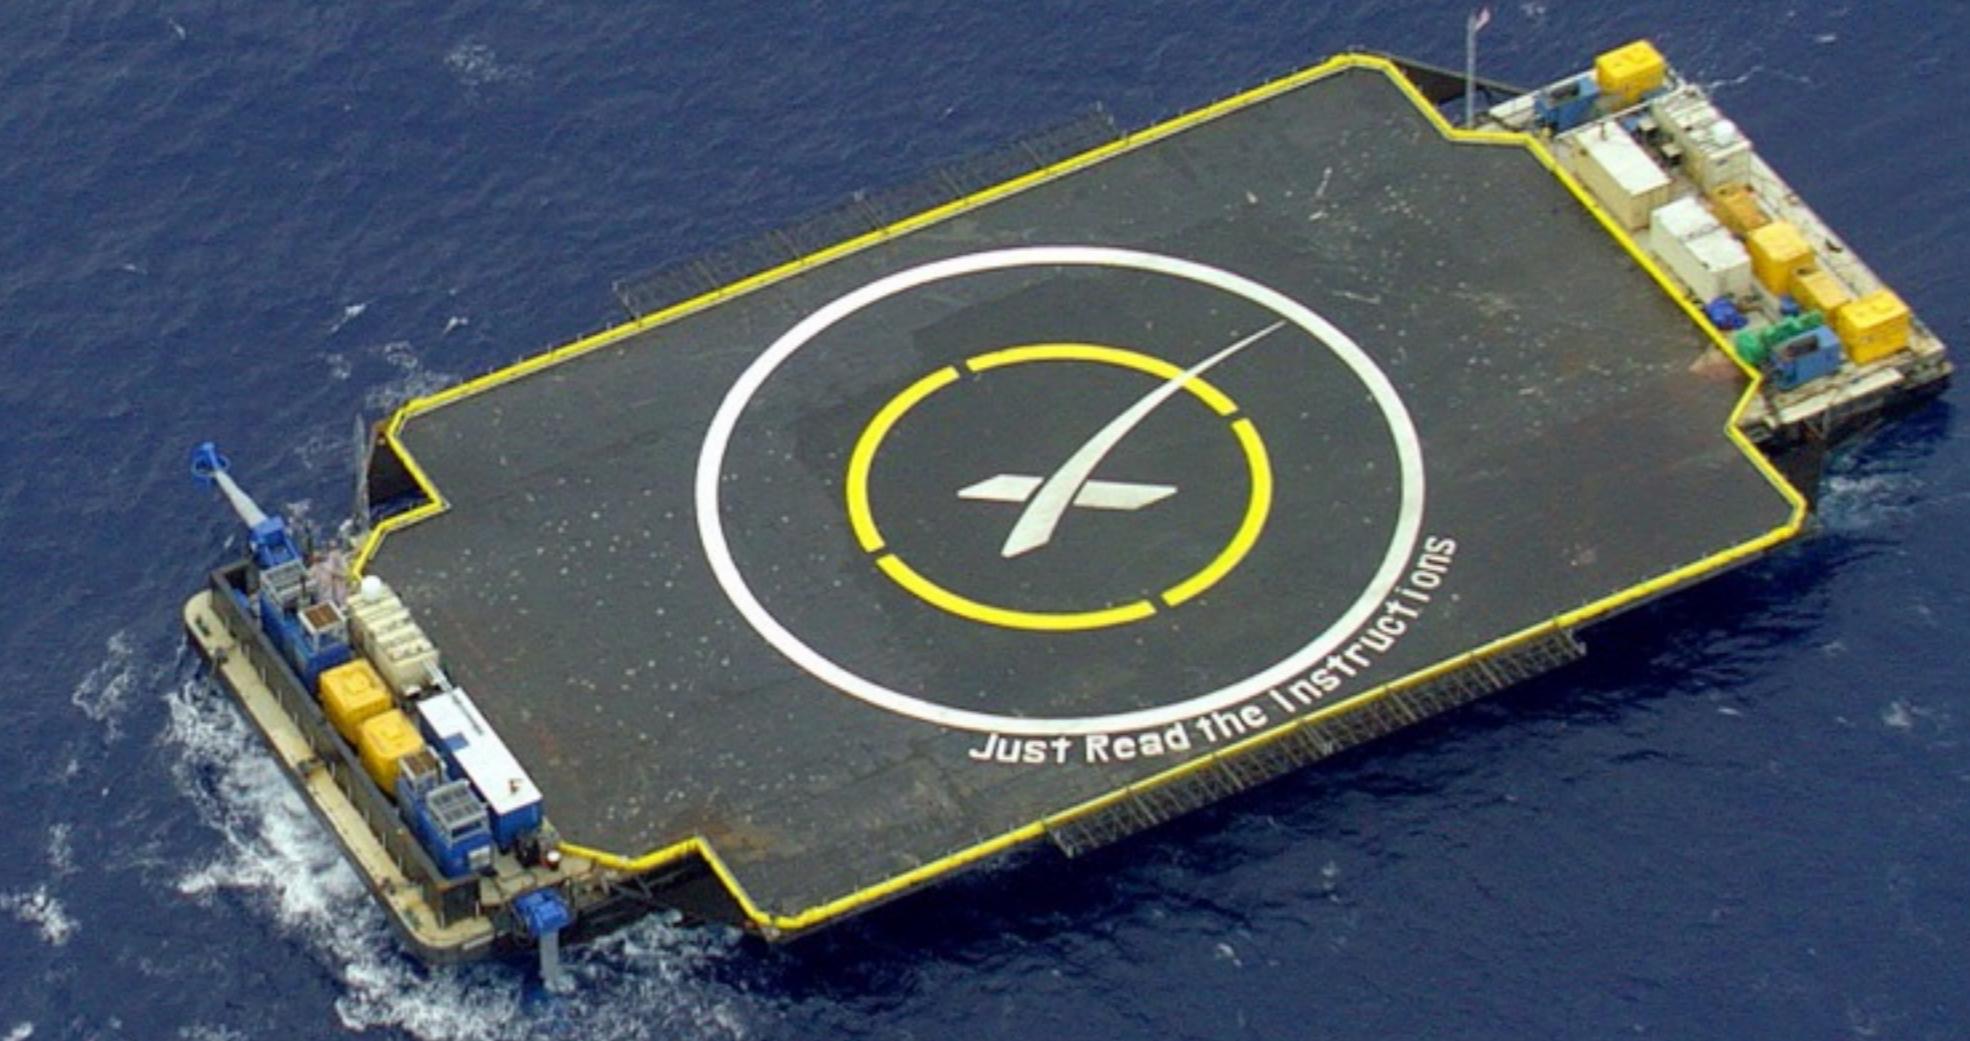
```
>>> hp.setrelheap() # Used as reference point for memory usage
>>> run_task("tasks.your.MemoryHungryTask", {"a": 1, "b": 2})
>>> h = hp.heap()
```

```
>>> h
Partition of a set of 347 objects. Total size = 61320 bytes.
Index  Count   %     Size   % Cumulative   % Kind (class / dict of class)
    0    215   62    18920   31      18920  31 __builtin__.weakref
    1      4   1     8800   14      27720  45 dict of mongokit.document.DocumentProperties
    2     17   5     8328   14      36048  59 list
    3      4   1     5792    9      41840  68 mongokit.helpers.DotCollapsedDict
    4      8   2     4544    7      46384  76 dict (no owner)
    5      4   1     3616    6      50000  82 mongokit.document.DocumentProperties
    6      5   1     1160    2      51160  83 __builtin__.set
    7      4   1     1120    2      52280  85 dict of mongokit.helpers.DotCollapsedDict
    8      1   0     1048    2      53328  87 dict of 0x279e1f0
    9      1   0     1048    2      54376  89 dict of 0x2905040
<18 more rows. Type e.g. '_more' to view.>
```

Misc tools

- ▶ Scheduler
- ▶ Command-line runner, e.g. `mrq-run tasks.myTask {"a": 1}`
- ▶ Autoscaling
- ▶ Profiler

CONSISTENCY



Consistency guarantees

- ▶ At least once vs. At most once vs. Exactly once
- ▶ Ordering
- ▶ Critical operations:
 - ▶ Queueing
 - ▶ Marking tasks as started
 - ▶ Timeouts & retries

Types of brokers

- ▶ Specialized message queues (RabbitMQ, SNS, Kafka, ...)
 - ▶ Performance, complexity, poor visibility
- ▶ In-memory data stores (Redis, ...)
 - ▶ Performance, simplicity, harder to scale
- ▶ Regular databases (MongoDB, PostgreSQL, ...)
 - ▶ Often enough for the job!

At the heart of the broker

- ▶ Atomic update from "queued" to "started"
- ▶ MRQ with MongoDB broker: `find_one_and_update()`
- ▶ MRQ with Redis broker: Pushback in a ZSET

Queue type	Regular	Raw	Raw with <code>no_storage</code> config
Storage for queued jobs	MongoDB	Redis	Redis
Storage for started & success jobs	MongoDB	MongoDB	None
Performance	+	++	+++
Visibility in the dashboard	Full	After start	Job counts & failed jobs
Safety	+++	++	+

ZSETs in Redis

- ▶ Sorted sets with $O(\log(N))$ scalability
- ▶ set/get by key, order by key, lookups by key or value
- ▶ Very interesting properties for task queues: Unicity, Ordering, Atomicity of updates, Performance
- ▶ MRQ's "Pushback" model:
 - ▶ Queue with key=timestamp
 - ▶ Unqueue by fetching key range & setting new keys in the future
 - ▶ After completion the task adjusts or removes the key

```
def redis_zaddbyscore():
    """ Increments multiple keys in a sorted set & returns them """

    return context.connections.redis.register_script("""
local zset = KEYS[1]
local min = ARGV[1]
local max = ARGV[2]
local offset = ARGV[3]
local count = ARGV[4]
local score = ARGV[5]

local data = redis.call('zrangebyscore', zset, min, max, 'LIMIT', offset, count)
for i, member in pairs(data) do
    redis.call('zadd', zset, score, member)
end

return data
""")
```

Consistency guarantees

- ▶ Must be thought of for the whole system, not just the broker!
- ▶ Brokers can be misused or misconfigured
- ▶ The workers can drop tasks if they want to ;-)
- ▶ Consistency starts at queueing time!



TIME TO CHOOSE!

Think hard about what you need

- ▶ Will your taskqueue be the foundation of your architecture, or is it just a side project?
- ▶ What performance do you need? (IO vs. CPU, latency, ...)
- ▶ What level of visibility and control do you need on queued & running tasks?
- ▶ Can workers terminate abruptly? Lots of design consequences!
- ▶ What language interop do you need?

And then all the usual questions...

- ▶ Is it supported by a lively community?
- ▶ License
- ▶ Documentation
- ▶ Future plans

Which one to pick?

- ▶ **Celery:** High performance, large community, very complex, major upgrades painful
- ▶ **RQ:** Extremely simple to understand, low performance
- ▶ **MRQ:** Adjust task visibility vs. performance, simple to understand, 1.0 soon
- ▶ Lots of other valid options! Just be sure to ask yourself the right questions ;-)



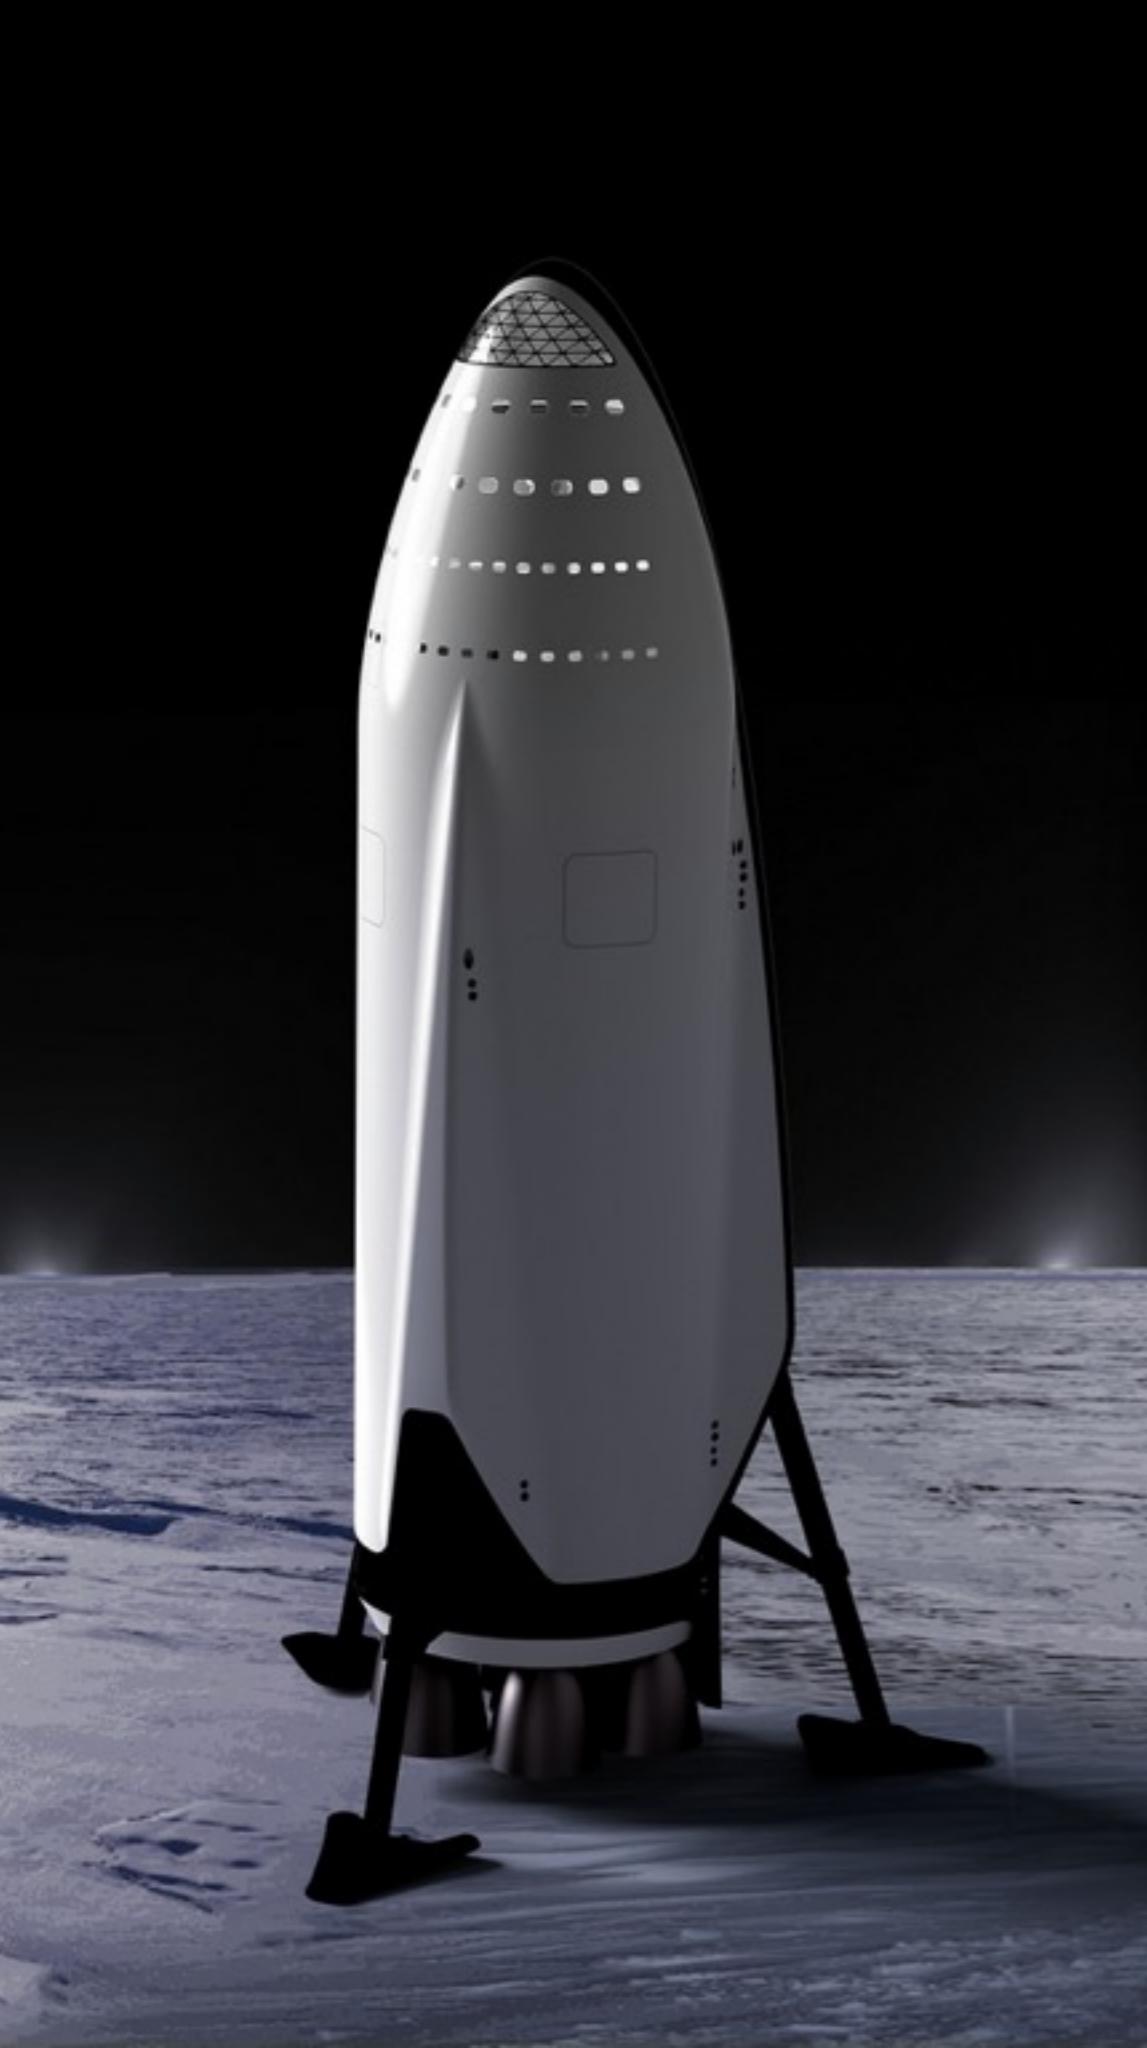
REMINDER

BE GRATEFUL FOR
THE OSS YOU USE!



PricingAssistant

Hiring Pythonistas!



THANKS!
—
QUESTIONS?